

Report on Logistic Regression for Diabetes Prediction

1. Introduction

Diabetes is a chronic disease that affects millions worldwide. Early detection and prediction play a crucial role in preventive healthcare. In this project, we use a medical dataset to build and evaluate a logistic regression model for predicting the likelihood of diabetes.

The dataset is a processed version of the Pima Indians Diabetes Dataset, with the following features:

- times_pregnant - Number of times the patient has been pregnant
- glucose_conc - Plasma glucose concentration
- Diastolic_BP - Diastolic blood pressure
- Triceps_thk - Triceps skin fold thickness
- 2_hr_insulin - 2-hour serum insulin level
- BMI - Body mass index
- Pedigree - Diabetes pedigree function (family history)
- Age - Patient's age
- Diabetes (target variable) - Encoded as 1 (diabetic) and 0 (non-diabetic)

The target variable was originally encoded as -1 and 1, which caused issues with binary classification metrics. We resolved this by remapping -1 to 0.

2. Exploratory Data Analysis (EDA)

Steps Performed:

1. Data Loading: The dataset was loaded into pandas from './Data/proc_pima_2_withheader.csv_preprocessed.csv'.
2. Missing Values: No missing values were present.
3. Statistical Summary: We inspected mean, median, standard deviation, and ranges to identify scale differences.
4. Feature Scaling: Features were standardized using StandardScaler.
5. Target Distribution: Balanced representation ensured fair evaluation.
6. Feature Engineering: Final dataset consisted of scaled features + binary target.

Exploratory Data Analysis (EDA)

Summary statistics

```
print("Summary Statistics:")  
print(df.describe())
```

[4]

... Summary Statistics:

	Diabetes	times_pregnant	glucose_conc	Diastolic_BP	Triceps_thk
count	394.000000	394.000000	394.000000	394.000000	394.000000
mean	-0.335025	3.302030	122.380711	70.675127	29.137056
std	0.943407	3.208235	31.422626	12.465567	10.503919
min	-1.000000	0.000000	0.000000	24.000000	7.000000
25%	-1.000000	1.000000	99.000000	62.000000	21.000000
50%	-1.000000	2.000000	119.000000	70.000000	29.000000
75%	1.000000	5.000000	143.000000	78.000000	36.750000
max	1.000000	17.000000	198.000000	110.000000	63.000000

	2_hr_insulin	BMI	Pedigree	Age
count	394.000000	394.000000	394.000000	394.000000
mean	155.322335	33.073858	0.522741	30.888325
std	118.987181	7.015055	0.344833	10.232549
min	0.000000	18.200000	0.085000	21.000000
25%	76.000000	28.400000	0.270250	23.000000
50%	125.000000	33.200000	0.449500	27.000000
75%	190.000000	37.075000	0.685750	36.000000
max	846.000000	67.100000	2.420000	81.000000

3. Logistic Regression Implementation

We implemented logistic regression using gradient descent optimization from scratch. The main steps were:

- Sigmoid Function: To map linear combinations of inputs to probabilities.
- Loss Function: Binary cross-entropy loss was minimized.
- Gradient Descent: Used to iteratively update weights until convergence.
- Thresholding: Predictions converted to 0 or 1 based on a 0.5 cutoff.

The scratch implementation was compared against scikit-learn's LogisticRegression for validation.

```

class LogisticRegressionScratch:
    def __init__(self, lr=0.01, epochs=1000):
        self.lr = lr
        self.epochs = epochs
        self.weights = None
        self.bias = None

    def sigmoid(self, z):
        return 1 / (1 + np.exp(-z))

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.weights = np.zeros(n_features)
        self.bias = 0

        for _ in range(self.epochs):
            linear_model = np.dot(X, self.weights) + self.bias
            y_pred = self.sigmoid(linear_model)

            dw = (1 / n_samples) * np.dot(X.T, (y_pred - y))
            db = (1 / n_samples) * np.sum(y_pred - y)

            self.weights -= self.lr * dw
            self.bias -= self.lr * db

    def predict_proba(self, X):
        linear_model = np.dot(X, self.weights) + self.bias
        return self.sigmoid(linear_model)

    def predict(self, X):
        y_pred_prob = self.predict_proba(X)
        return np.where(y_pred_prob >= 0.5, 1, 0)

```

4. Model Evaluation

We split the dataset into 80% training and 20% testing. Both implementations were evaluated on accuracy, precision, recall and F1 score

Results:

[Scratch Logistic Regression]


- Accuracy: 0.8228
- Precision: 0.7778
- Recall: 0.7241
- F1 Score: 0.7500

[Scikit-learn Logistic Regression]

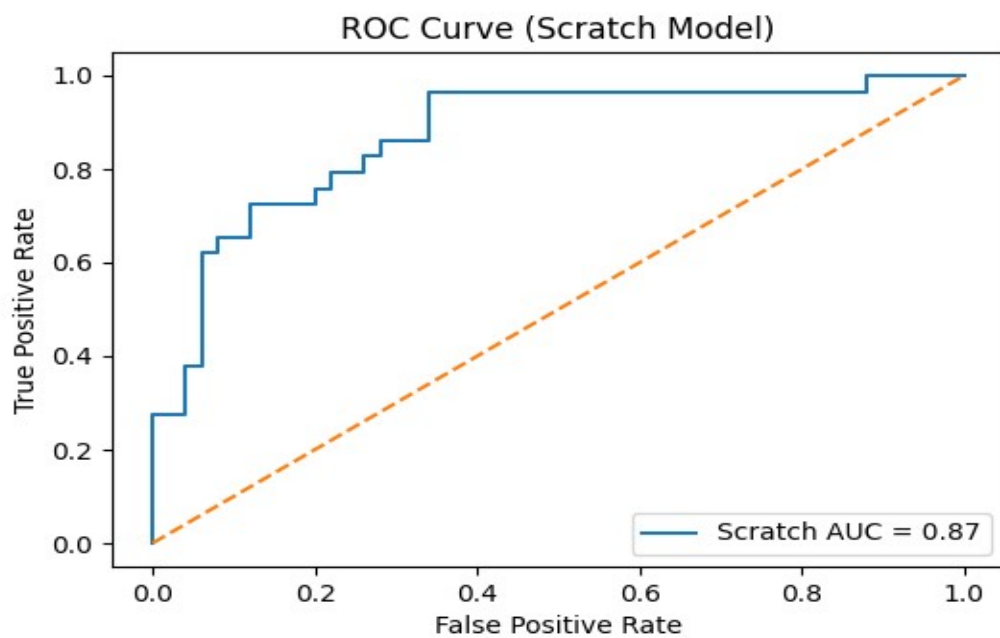
- Accuracy: 0.8228
- Precision: 0.7778
- Recall: 0.7241
- F1 Score: 0.7500

The results were identical, confirming the correctness of our scratch implementation.


- Confusion Matrix



- ROC Curve



- Precision-Recall Curve



5. Conclusion

The project demonstrated that logistic regression can be successfully implemented from scratch and achieve the same performance as scikit-learn's library implementation.

Key takeaways:

- Proper preprocessing (feature scaling, binary target encoding) is essential.
- Gradient descent optimization produced stable convergence.
- Evaluation metrics showed good predictive performance (~82% accuracy).
- Confusion matrix and ROC/PR curves confirmed balanced trade-offs between precision and recall.

This project reinforces logistic regression as a robust baseline classifier for binary medical datasets and provides a foundation for experimenting with more advanced machine learning algorithms.